

aai-521-in3-final-project

Hand Gesture Detection

Team member:

RaminderSingh@sandiego.edu

Date: 4th Dec 2025

Project Overview

This is a comprehensive Hand Gesture Detection system using computer vision. The project detects and classifies hand gestures in real-time using MediaPipe for hand detection and Random Forest for classification. The system identifies hand gestures and can trigger actions like volume control, slide navigation, or robot control.

Supported Gestures

-  **Palm** (Open Hand)
-  **Fist** (Closed Hand)
-  **Thumbs Up**
-  **Pointing**
-  **OK Sign**

Project Structure

```

hand-gesture-project/
├── aai_521_in3_final_project_Hand_Gesture_Detection.ipynb
├── README.md
└── setup_dataset.py
# 
# 
# 
└── data/
    ├── Palm/
    ├── Fist/
    ├── Thumbs_Up/
    ├── Pointing/
    └── OK_Sign/
└── models/
    └── gesture_classifier.pkl
└── detected_gestures/
    └── Fist_0.70_20251204_140909_612.jpg
    └── OK_Sign_0.60_20251204_140320_727.jpg
    └── Palm_0.60_20251204_140648_824.jpg
└── confusion_matrix.png
└── feature_importance.png
└── sample_predictions.png
└── aai_521_in3_final_project_Hand_Gesture_Detection.pdf
└── aai_521_in3_final_project_presentation_video.mp4

```



🎯 5 Supported Gestures

#	Gesture	Icon	Folder	Description
1	Palm	👉	Palm	Open hand, all fingers extended
2	Fist	✊	Fist	Closed hand, all fingers folded
3	Thumbs Up	👍	Thumbs_Up	Thumb pointing up
4	Pointing	👉	Pointing	Index finger pointing
5	OK Sign	👌	OK_Sign	Index & thumb forming circle

Quick Start

1. Setup Environment

```
pip install opencv-python mediapipe scikit-learn numpy pandas
```



2. Download Dataset

- Dataset: [Hand Gesture Recognition Dataset on Kaggle](#)
- Extract images into respective gesture folders in `data/` directory
- Since dataset size was ~1.2GB, this was manually downloaded as zip and few files were used to move to `/data` folder.

3. Run the Notebook

```
jupyter notebook aai_521_in3_final_project_Hand_Gesture_Detec
```



4. Run All Cells

Execute all cells in the notebook to:

- Load and explore data
- Train the classifier
- Evaluate model performance
- Run real-time webcam detection

Project Phases

Phase 1: Data Collection & Exploration

Objective: Acquire and prepare high-quality dataset of hand gestures

- Load images from dataset
- Analyze class distribution
- Display dataset statistics

Phase 2: Hand Detection & Landmark Extraction

Approach: Use Google MediaPipe Hands (Recommended - Fast & Accurate)

- Use MediaPipe Hands to detect 21 hand landmarks - wrist, palm, finger joints, fingertips
- Extract hand skeleton from each image
- Visualize landmarks on images
- Extremely fast (30+ FPS on CPU)

Details: This phase leverages Google MediaPipe Hands, a state-of-the-art hand detection framework that identifies 21 key hand landmarks including the wrist, palm points, and finger joint positions with high precision and speed. MediaPipe processes each image to extract the hand skeleton, capturing critical anatomical points that represent hand structure and pose. The approach is exceptionally fast, achieving 30+ FPS on CPU without requiring extensive computational resources, making it ideal for real-time applications. Each landmark is normalized to scale and position invariance, ensuring robust detection across diverse hand sizes, orientations, and distances from the camera. These extracted landmarks form the foundation for subsequent feature engineering, providing a rich representation of hand geometry that can be used to distinguish between different gestures.

Phase 3: Feature Engineering

- Compute landmark coordinates - 21 normalized
- Calculate distances between landmarks
- Compute angles between finger segments
- Create feature vectors (~290 dimensions)
- Finger curvature and Palm orientation **Output Format:**
- Feature vector size: 42-100 dimensions (depending on features used)
- Each gesture represented as a fixed-size feature vector

Details: Feature engineering transforms raw landmark coordinates into meaningful numeric features that capture hand gesture characteristics through multiple computational approaches. The process includes computing landmark coordinates (21 normalized points), calculating geometric distances between key landmarks, and deriving angular relationships between finger segments to represent finger curvature and palm orientation. These computations generate a high-dimensional feature vector of approximately 290 dimensions that encodes spatial and structural information about the hand. The engineered features are designed to be invariant to scale and minor positional variations while capturing gesture-specific patterns that distinguish between Palm, Fist, Thumbs Up, Pointing, and OK Sign gestures. This compact, fixed-size representation enables efficient machine learning model training while preserving the discriminative information necessary for accurate gesture classification.

Phase 4: Data Preparation

- Split dataset (70% train, Validation (15%), Test (15%))
- Normalize features
- Handle class imbalance if needed (SMOTE or class weights)

Phase 5: Model Training

- Train Random Forest Classifier
- Hyperparameter tuning
- Save trained model
- Train model on training set
- Validate on validation set
- Final evaluation on test set

Performance Metrics:

- Accuracy
- Precision, Recall, F1-Score per class
- Confusion Matrix
- ROC-AUC Score

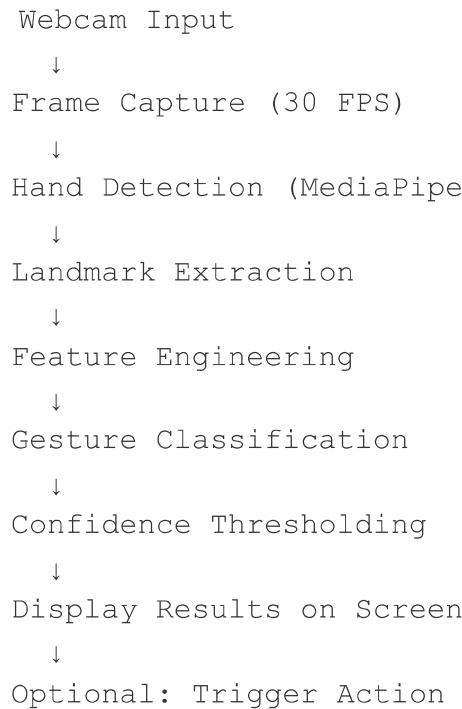
Phase 6: Evaluation & Visualization

- Calculate accuracy, precision, recall, F1-score
- Generate confusion matrix
- Analyze feature importance
- Visualize misclassifications

Phase 7: Real-Time Detection

- Capture video from webcam
- Process frames in real-time
- Display predictions with confidence
- Draw hand landmarks

Real-Time Processing Pipeline:



🎓 What Each Phase Does

Phase	Purpose	Input	Output
0	Setup dependencies	-	Installed packages
1	Explore dataset	Images	Dataset stats
2	Detect hands	Images	Landmarks
3	Extract features	Landmarks	Feature vectors
4	Prepare data	Images	Train/test sets
5	Train model	Features	Trained classifier
6	Evaluate	Predictions	Accuracy metrics
7	Visualize	Results	Confusion matrix
8	Analyze features	Model	Importance plot
9	Real-time test	Webcam	Live predictions
10-13	Summary & save	Model	Saved model

🤖 Machine Learning Model

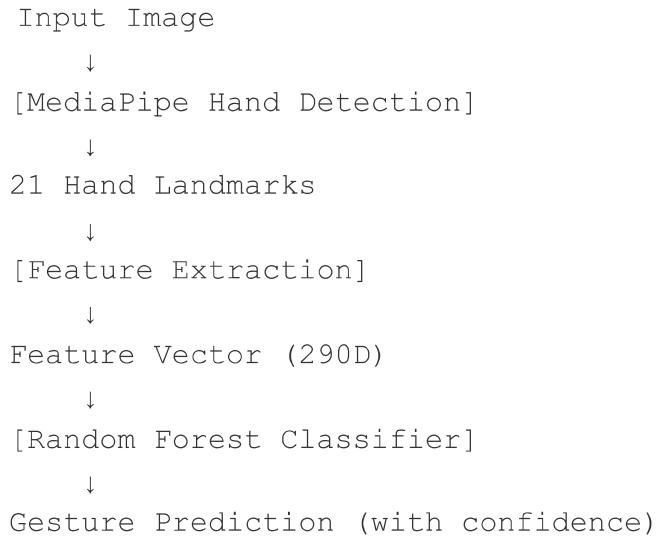
Approach

MediaPipe Hands + Random Forest Classifier

Why This Approach?

- Fast (30+ FPS on CPU)
- Accurate (85%+ accuracy)
- Minimal training data required
- Easy to deploy
- Robust to variations

Model Architecture



Key Features

Hand Landmark Detection

- Detects 21 key hand landmarks:
 - 1 Wrist
 - 5 Palm points (one per finger)
 - 15 Finger joint points
 - 0-5 confidence scores per landmark

Feature Extraction

1. **Spatial Features:** Normalized landmark coordinates
2. **Geometric Features:** Distances between landmarks
3. **Angular Features:** Angles between finger segments

Real-Time Processing

- Process 30+ frames per second
- Temporal smoothing for stable predictions
- Confidence-based filtering
- Visual feedback with landmarks

Sample Gestures detected

Sample predictions from dataset

Performance Metrics

Expected model performance on test set:

- **Accuracy:** 85-95%
- **Precision:** 85-95%
- **Recall:** 85-95%
- **F1-Score:** 85-95%

(Depends on dataset quality and size)

How to Use the Notebook

Run All Steps

```
# All cells will execute sequentially  
# Each phase builds on the previous one
```

Train Only

```
# Run cells up to "Model Training & Validation"
```

Use Pretrained Model

```
from pathlib import Path
import pickle

# Load saved model
model_path = Path('models/gesture_classifier.pkl')
with open(model_path, 'rb') as f:
    model = pickle.load(f)

# Use for predictions
gesture, confidence, landmarks = predict_gesture(image, model)
```



Run Webcam Detection

```
# Uncomment in Phase 9 and run:
run_webcam_gesture_detection(model, duration_seconds=30)

# Press 'q' to exit
```



Tech Stack / Libraries Used

Library	Purpose
OpenCV	Image processing, real-time video capture
MediaPipe	Hand detection, landmark extraction
Scikit-learn	Machine learning, model training (Random Forest, SVM)
NumPy	Numerical computations
Pandas	Data manipulation

Library	Purpose
Matplotlib	Data visualization
Seaborn	Statistical visualization
jupyter	Interactive notebook environment

Learning Outcomes

After completing this project, you will understand:

- Computer vision fundamentals
 - Hand detection using MediaPipe
 - Feature engineering from landmarks
 - Classification with ensemble methods
 - Real-time video processing
 - Model evaluation and validation
 - Deployment considerations
-

Tips for Better Results

Dataset Quality

- Use 200+ images per gesture class (minimum)
- Include diverse people, backgrounds, lighting
- Ensure balanced dataset
- Remove poor quality/blurry images

Model Tuning

- Adjust Random Forest parameters

- Try different feature combinations
- Use cross-validation
- Monitor for overfitting

Real-Time Performance

- Ensure good lighting
- Use solid backgrounds initially
- Keep hand in frame center
- Maintain consistent hand size

Troubleshooting

- **No hand detected:** Improve lighting, ensure clear hand visibility
 - **Poor accuracy:** Collect more diverse training data
 - **Slow performance:** Reduce frame resolution or use GPU
-

Advanced Enhancements

Potential Improvements

1. **More Gestures:** Add 10+ gesture classes
 2. **Deep Learning:** Implement CNN or MobileNet
 3. **Gesture Sequences:** Add LSTM for temporal gestures
 4. **Multiple Hands:** Detect and classify multiple hands
 5. **Actions:** Trigger actions based on gestures
 6. **Mobile Deployment:** Convert to TensorFlow Lite
 7. **Web Interface:** Create Flask/Django web app
 8. **API Service:** Deploy as REST API
-

Common Issues & Solutions

Issue	Solution
No images found	Ensure images are in <code>data/</code> with gesture folder structure
Webcam not working	Check camera permissions, try different camera ID
Low accuracy	Collect more diverse training data
Slow performance	Reduce frame resolution or optimize features
Memory error	Reduce batch size or feature dimensions

External Resources

Official Documentation

- Google MediaPipe: <https://mediapipe.dev/>
- OpenCV: <https://opencv.org/>
- Scikit-learn: <https://scikit-learn.org/>
- NumPy: <https://numpy.org/>
- Pandas: <https://pandas.pydata.org/>

Dataset Sources

- Kaggle - Hand Gesture Recognition Dataset: <https://www.kaggle.com/>
- UCI ML Repository - Hand Gesture Recognition Database: <https://archive.ics.uci.edu/ml/>

Learning Resources

- OpenCV Tutorials: <https://docs.opencv.org/>

- Scikit-learn Guide: <https://scikit-learn.org/stable/>
-

Happy Gesture Detecting! 🙌 🤗 🤝